

Mapping Tinder's Safety Feature Claims to Problem Areas and Risks

J Shelly¹ - Northeastern University, Khoury College of Computer Science MS Student, Thaddeus Kopczynski² - Northeastern University, Khoury College of Computer Science BS Student

Abstract

Dating apps expose users to risks at every stage of use, from fake profiles and harassing messages to harm during in-person meetings, and dating platforms describe a range of safety features meant to address them. These descriptions shape what users understand about the protection a feature offers. We examine how Tinder, the most widely used dating app in the United States, describes its safety features across its help center, newsroom, and policy site. We extract passages that describe safety features, group the features into problem areas, the contexts of use in which risks arise, and map each feature to the risks it claims to address. To classify risks, we build on frameworks for children's online safety: the 4Cs (Content, Contact, Conduct, and Contract), the 5Cs, which add circulation, and the 6Cs, which add compulsivity. We adapt these frameworks for adult users and extend them to the 7Cs by adding Community risk, which covers risks that arise once users meet offline or encounter each other within existing social networks. We identify problem areas that follow the arc of dating-app use, from establishing who a user is to meeting in person. Most of Tinder's described features address Contact risk, followed by Community risk, while Content and Compulsivity risks are addressed by few features. A preliminary reflexive thematic analysis further shows that Tinder presents effectiveness claims without context, places responsibility for safety on users while leaving its own enforcement criteria undefined, and varies feature availability by region and device, often without explanation. Our contributions are the 7Cs risk taxonomy, a framework for mapping dating-app safety features to problem areas, a map of where across those areas each type of risk is mitigated, and an account of how Tinder communicates about its safety features.

Keywords: dating apps, online safety, safety features, platform communication, risk taxonomy

1 Introduction

Dating apps have become a common way to meet new people. Three in ten U.S. adults report having used a dating site or app, and Tinder is the most widely used dating platform in the United States [Vogels and McClain 2023]. Their users face risks at every stage of use, from fake profiles and harassing messages to romance scams and sexual violence during in-person meetings [Bilz et al. 2023; Filice et al. 2022; Gillett 2018; O'Brien et al. 2025]. To address these risks, dating platforms describe a range of safety features, from identity verification to tools for sharing date plans with friends [Dietzel et al. 2026; O'Brien et al. 2025].

The safety practices of dating platforms have drawn regulatory and public scrutiny. In 2019, the U.S. Federal Trade Commission (FTC) sued Match Group, Tinder's parent company, alleging that its Match.com service used messages from accounts it had flagged as likely fraudulent to sell subscriptions, and that it protected paying subscribers, but not free users, from some of those accounts [CBS News 2019]. In 2025, Match Group agreed to pay \$14 million to settle the case [Morris 2025]. That same year,

an investigative reporting project found that Match Group had not released a transparency report on harm that it had promised in 2020 [The Dating Apps Reporting Project 2025].

Research on online safety commonly organizes risks using the 4Cs framework, which classifies online risks to children as content, contact, conduct, and contract risks [Livingstone and Stoilova 2021]. Ma et al. [2026] extended the 4Cs to the 5Cs by adding circulation risk, the spread of content beyond its intended audience, and Matsumoto et al. [2026] extended the 5Cs to the 6Cs by adding compulsivity risk, which covers excessive or mistimed use. All three frameworks were developed for children's use of social media. Applying them to dating apps requires adapting their definitions for adult users and accounting for risks that arise once users meet offline or encounter each other within existing social networks.

Recent studies have examined dating platforms' own safety materials. Dietzel et al. [2026] analyzed terms of service, privacy policies, community guidelines, and safety tips from 30 dating apps and found that they individualize responsibility for safety, rely on peer and law-enforcement policing, and promote surveillance-based features such as identity verification. O'Brien et al. [2025] reviewed formal and informal documents from five popular dating platforms in the United Kingdom, including Tinder, and examined how the platforms characterize harms, respond to violations, and describe their safeguarding tools; they found that many safeguards shift the burden of safety onto users.

How platforms describe their safety features matters, because users rely on these descriptions to understand and adopt the features [Ma et al. 2026]. Independent audits of social media platforms have found that many advertised safety features do not work as described [Béjar et al. 2025; Matsumoto et al. 2026]. Establishing what a dating platform claims its safety features do is a necessary first step toward evaluating them.

In this paper, we analyze how Tinder describes its safety features in three public channels: its help center, newsroom, and policy site. Following Ma et al. [2026], we use the concept of *problem areas*, the contexts of use in which risks can arise, to connect the features a platform describes to the risks those features claim to address. We ask two research questions:

- **RQ1:** What problem areas characterize the contexts in which risks arise for Tinder users, and where does Tinder claim to apply safety features?
- **RQ2:** How does Tinder communicate about its safety features across its help center, newsroom, and policy site?

To answer these questions, we collected pages from the three channels and extracted passages that describe safety features. We standardized the described features, grouped them into problem areas, and mapped each feature to the risks it claims to address. For this mapping, we adapt the 6Cs [Matsumoto et al. 2026] for adult users and extend them to the 7Cs by adding Community risk, which covers risks that arise once users meet offline or encounter each other within existing social networks. For RQ2, we conducted a preliminary reflexive thematic analysis [Braun and Clarke 2019], first coding the extracted passages article by article and then comparing each feature's descriptions across the three channels to identify common themes.

We identify problem areas that follow the arc of dating-app use, from establishing who a user is to meeting in person. Two of them, Identity & Authenticity and In-Person / Offline Safety, have no

counterpart in the problem-area framework Ma et al. [2026] developed for youth-oriented social media. Mapping features to the 7Cs shows that Tinder's described features most often address contact risk, followed by Community risk, which concentrates around in-person meetings; content and compulsivity risks are addressed by only a few features. Our thematic analysis identifies three patterns in how Tinder communicates about these features: effectiveness claims presented without context, responsibility for safety placed on users while enforcement criteria remain undefined, and feature availability that varies by region, device, or partner, often without explanation.

This paper makes four contributions. First, it extends the 6Cs to the 7Cs by adding Community risk, which captures risks that arise once users meet in person or encounter each other within existing social networks. Second, it provides a framework for mapping a dating platform's safety features to problem areas, adapting the approach of Ma et al. [2026] to dating apps. Third, using the 7Cs and this framework, it maps where across Tinder's problem areas each type of risk is mitigated, which researchers, regulators, and platforms can use to evaluate a platform's safety coverage. Fourth, it identifies three patterns in how Tinder communicates about its safety features across its help center, newsroom, and policy site.

The remainder of this paper is organized as follows. Section 2 reviews research on dating-app risks, the risk frameworks we build on, safety features on dating apps, how platforms communicate about safety, and the relationship between monetization and safety. Section 3 describes our data collection and analysis. Section 4 presents the problem areas, the 7Cs, where each risk is mitigated, and how Tinder communicates about its safety features, and Section 5 discusses their implications. Section 6 covers limitations and future directions, and Section 7 concludes.

2 Background

We review research on the risks dating apps pose, the risk frameworks this study builds on, safety features on dating apps, how platforms communicate about safety, and the relationship between monetization and safety.

2.1 Dating Apps and Risk

Dating apps support friendships, relationships, and other connections, but they can also be sites of risk and harm [Dietzel et al. 2026]. Gender-based and sexual violence are frequently facilitated through dating apps [Filice et al. 2022], and sexually aggressive messages and other intimate intrusions have become common enough to go unaddressed [Gillett 2018]. Women are targeted more often than men by romance scams [Bilz et al. 2023], racialized women report fetishization [Banks et al. 2024], and LGBTQ+ users report discrimination and threatening behavior [Albury et al. 2021]. Women, sexual and gender minorities, and people of color face disproportionate risk [Filice et al. 2022].

Dating-app risks also differ from those of other platforms. Social media and gaming platforms support online interaction as an end in itself, whereas dating platforms use online interaction to arrange offline meetings [O'Brien et al. 2025]. Physical distance and relative anonymity can enable deceptive self-presentation, and access to personal information can enable stalking and harassment [O'Brien et al. 2025].

Users also develop their own strategies to manage these risks, such as the safety work women perform on Tinder [Gillett 2023]. Our study examines how a platform describes the features it offers to reduce these risks.

2.2 Risk Frameworks

The 4Cs framework classifies online risks to children into four types [Livingstone and Stoilova 2021]. Content risks involve exposure to harmful material, contact risks involve harmful interactions initiated by adults, conduct risks involve children's own behavior, and contract risks involve commercial and data exploitation. HCI and CSCW researchers have used the framework to evaluate safety mechanisms such as parental control apps [Wang et al. 2021]. Ma et al. [2026] added circulation risk, the spread of content beyond its intended audience, to form the 5Cs. Matsumoto et al. [2026] added compulsivity risk, which covers excessive or mistimed use, to form the 6Cs.

These frameworks were developed for children's use of social media. Applying them to a dating app requires two changes. First, because dating-app users are adults, we define contact and conduct risks by who initiates an interaction rather than by the age of the people involved. Second, none of the frameworks was developed for a product whose intended outcome is an offline meeting. We therefore add Community risk, which covers risks that arise once users meet in person or encounter each other within existing social networks.

2.3 Safety Features on Dating Apps

Dating platforms offer a range of built-in safety features. These include blocking and reporting tools, screenshot blocking, location safeguards for LGBTQ+ users traveling to countries with laws that target them, and prompts that ask senders to reconsider potentially harmful messages [O'Brien et al. 2025]. They also include photo and identity verification, third-party background checks, tools for sharing date plans with friends, emergency-service integrations, and automated message filters [Dietzel et al. 2026].

Researchers have questioned how well these features protect users. Reactive features such as blocking place responsibility on users after harm has occurred [Aljasim and Zytka 2023], and automated detection raises concerns about bias and privacy [O'Brien et al. 2025]. Data-intensive verification and policing features have been criticized for equating surveillance with safety [Stardust et al. 2023]. Design research has explored alternatives, such as participatory design of risk-detection tools intended to help women assess potential harms [Datey and Zytka 2024].

We organize Tinder's safety features using the problem-area approach that Ma et al. [2026] developed for social media platforms used by youth, mapping each feature to the context in which it applies and the risks it claims to address.

2.4 How Platforms Communicate About Safety

Platforms announce and explain safety features through public channels such as newsrooms and help centers. Ma et al. [2026] analyzed such communication from YouTube, TikTok, Meta, and Snapchat and identified three problematic practices: claims about implementation that did not match feature availability, unclear explanations of how features operate, and a lack of evidence that features reduce risk.

Studies of dating platforms report related patterns. Dietzel et al. [2026] found that dating apps' safety materials often hold users responsible for their own safety and give little detail about what happens after a report. O'Brien et al. [2025] found that platforms rarely define the harms they prohibit, that some self-protection advice conflicts with platform rules, and that safeguards vary by region and by user group.

Independent audits test whether such claims hold. Béjar et al. [2025] tested Instagram's teen safety features with scenario-based test accounts and found that most were ineffective or no longer available. Matsumoto et al. [2026] audited child safety features on Instagram, Snapchat, TikTok, and YouTube and found that most failed. Our study follows Ma et al. [2026] in analyzing feature-level claims across a platform's public channels and extends this approach to a dating app.

2.5 Monetization and Safety

Dating platforms commonly combine free accounts with paid subscriptions; Match.com, for example, let users create profiles for free but required a subscription to reply to messages [CBS News 2019]. This structure can shape safety. The FTC's 2019 complaint alleged that Match.com protected paying subscribers, but not free users, from some accounts it suspected of fraud [CBS News 2019]. Dietzel et al. [2026] note that market logics can commodify safety features, with unequal consequences for marginalized users [Campaioli et al. 2025]. Some safety resources are also limited by region; BrightCheck background checks, which Tinder signposts, are available only in the United States [O'Brien et al. 2025]. We therefore record whether each feature Tinder describes is free or requires payment.

3 Methods

We adapt the approach of Ma et al. [2026], who analyzed social media platforms' public communication about youth safety features, to a single dating platform. Our analysis has two parts. For RQ1, we map the safety features Tinder describes to problem areas and to the risks those features claim to address. For RQ2, we conduct a thematic analysis of how Tinder describes these features, first article by article and then feature by feature across channels.

3.1 Platform and Sources

We focus on Tinder, the most widely used dating platform in the United States [Vogels and McClain 2023]. Studying a single platform in depth lets us compare how the same feature is described in different places.

We collected pages from three public channels through which Tinder communicates about safety (Table 1). Each channel plays a different role. The help center explains how features work, the newsroom announces and promotes them, and the policy site states the rules and safety resources that govern their use. Ma et al. [2026] analyzed newsroom press releases and help or safety center articles. Prior studies of dating apps focused on formal and informal governance documents such as terms of service, community guidelines, and safety tips [Dietzel et al. 2026; O'Brien et al. 2025]. Including all three channels lets us examine whether Tinder describes a feature consistently when it announces, explains, and codifies it.

Table 1. Sources and collection yield.

Source	Site	Pages collected	Collection date
Help center	help.tinder.com	133	September 3, 2026
Newsroom	tinderpressroom.com	168	September 9, 2026
Policy site	policies.tinder.com	39	September 24, 2026
Total		340	

3.2 Data Collection

Data on Tinder's public communications were collected from three sources: the Tinder Newsroom (press releases), the Tinder support page, and Tinder's policy site. For each, we searched for sitemaps to enumerate available pages. This was followed by per-page scraping to extract metadata such as publish dates and categories. Where sitemaps proved incomplete (as with the policy site), a recursive link-crawling approach was used instead. Data cleaning addressed several platform specific issues uncovered during validation: navigation and language-switcher links that had initially contaminated the newsroom's article set were filtered out, and mass-shared. The policy site also required handling of inconsistent URL localization schemes and server-side locale guessing to ensure English-language pages were correctly identified and deduplicated. This process yielded 168 press releases, 133 help page articles, and 39 policy articles.

3.3 Relevance Screening

We screened each page for relevance, adapting the inclusion criteria of Ma et al. [2026], who retained articles that mentioned youth, online safety, and a technical feature or implementation. Because Tinder is an adult platform, we dropped the youth criterion. A page was relevant if it addressed online safety and described a safety feature or how one is implemented, including how Tinder detects, reviews, and acts on violations of its policies.

3.4 Quotation Extraction and Feature Coding

From each relevant page, we extracted passages about safety features, ranging from a single sentence to a short paragraph, and kept them verbatim. We extracted two kinds of passages. The first describes a specific mechanism: a feature a user could point to, turn on or off, or be shown, or a platform process that acts on users' accounts or content, such as detecting, reviewing, or enforcing against violations. The second describes how Tinder talks about a specific feature, including its stated purpose, the commitments Tinder makes about it, and claims about its effectiveness. The features these passages refer to are mapped in Section 3.5, and both kinds of passages are analyzed in Section 3.6. General safety advice that did not refer to a specific feature was not extracted. For each passage, we recorded the feature it refers to, the party it affects, the entity that controls it, any sub-features, whether it requires a paid subscription, and its source channel (Appendix A). A passage that discusses several closely related features, such as unmatching and blocking, was tagged to each of them. In Section 4.4, we identify quoted passages by source channel and, where available, year of publication; each passage's URL is recorded in our extraction table.

Several Tinder features share an underlying process, such as the steps common to its photo and identity verification features. We extracted a shared process once, under the base feature, and recorded only how each related feature differs. We then standardized feature names in two stages, following Ma et al. [2026]. Each passage first received a code in Tinder's own terminology, and codes describing the same mechanism were then merged. For example, Long Press Reporting, which opens the reporting flow when a user presses and holds a message, was merged into Report Message as another way to reach the same mechanism.

3.5 Mapping Features to Problem Areas and Risks (RQ1)

A problem area is the context or moment of use in which a type of risk can arise. Problem areas link a risk taxonomy to the concrete features a platform describes [Ma et al. 2026]. Following Ma et al., we derived problem areas bottom-up by grouping features according to the area of user activity they affect.

We then mapped each feature to the risks it claims to address. We started from the 6Cs taxonomy [Matsumoto et al. 2026], adapted its definitions for adult users, and added Community risk to cover risks that arise once users meet offline or encounter each other within existing social networks (Section 4.2). A feature could address more than one risk. We coded a feature as addressing contact risk if it protects users from other users, for example by limiting who can see or reach them, verifying or screening other users, signaling who a user is and what they are looking for, or enabling reporting, blocking, or removal. We coded it as addressing conduct risk if it shapes or responds to a user's own behavior: guiding what users send or disclose about themselves, educating them about safe behavior, or responding to their own violations through warnings, suspensions, and appeals. We coded content risk when a feature addresses exposure to harmful or inappropriate material, and contract risk when it concerns how the platform or its partners collect, use, or share users' data. We coded Community risk when a feature concerns meeting in person, people a user already knows, or being located or identified offline.

Finally, following Ma et al. [2026], we assess coverage. For each risk type, we treat all features that address it as 100% and compute the share that falls in each problem area. The result shows where in the use of the platform each type of risk is mitigated. We also report the number of features in each problem area (Table 2). Safety Center, a hub that links to other features rather than a feature in its own right, has no problem area and is excluded from these counts.

3.6 Thematic Analysis of Feature Descriptions (RQ2)

To examine how Tinder communicates about its safety features, we conducted a preliminary reflexive thematic analysis [Braun and Clarke 2006, 2019]. Following Ma et al. [2026], we chose reflexive thematic analysis because our goal was to identify recurring patterns in how features are described, which requires interpreting the text rather than counting terms. The analysis proceeded in three steps.

First, we read each article in full and coded the passages extracted from it, assigning short codes that describe what each passage does. Five questions guided coding: which risk the passage names and how precisely; what it claims the feature does and whether it offers evidence; who it makes responsible for safety; which conditions or limits it states or omits, such as subscription tier, region, or trigger threshold; and whether it frames the feature in terms of safety or in terms of fun, convenience, or quality. Coding also drew on the analytic notes recorded for each passage during extraction (Appendix A).

Second, we grouped passages by feature and compared each feature's descriptions across the help center, newsroom, and policy site. For each feature, we noted which channels describe it and whether their descriptions agree, differ in framing, or appear in only one channel.

Third, we grouped the codes from both steps into themes, patterns of shared meaning that recur across features and channels. We report themes supported across multiple features; a pattern observed in only a few passages is noted as a candidate for future analysis (Section 6.2). A single researcher conducted the analysis in two rounds of coding.

4 Findings

We present the problem areas we identified on Tinder, the risk taxonomy we use to map Tinder's safety features to these areas, and where across the problem areas each type of risk is mitigated (RQ1). We then present three themes in how Tinder communicates about its safety features (RQ2).

4.1 Problem Areas on Tinder

Table 2 lists the problem areas we identified. Each describes a context of use in which risks can arise for Tinder users and in which Tinder describes safety features.

Table 2. Problem areas on Tinder and the safety features Tinder describes in each.

Problem area	Context of use	Risks that arise	Described safety features (n)
Identity & Authenticity	Establishing that a profile represents a real person as presented	Profiles that misrepresent who a user is, from catfishing to scams	Photo Verification, Photo Check, Face Check, ID Verification, Email & Phone Verification, Background Check, Bio Guidance, Photo Enhance, Pronouns, Relationship Goals, Relationship Types (11)
Contact & Visibility Control	Controlling who can see, reach, or stay connected to a user	Unwanted or persistent contact; being found by people a user wants to avoid	Unmatch, Block Profile, Block Contacts, Report Profile, Report Message, Mutual Matching, Photo Verified Chat, Incognito, Disable Discovery, Search Restrictions, Don't Show My Age, Don't Show My Distance, Gender & Orientation Display,

Problem area	Context of use	Risks that arise	Described safety features (n)
			Automated Decision-Making, Cross-Platform Ban Notification (15)
In-App Communication	Exchanging messages, video, and contact details before meeting	Harassment, unwanted sexual messages, and pressure to move a conversation off the platform	Are You Sure?, Does This Bother You?, Contact Exchange Alert, Dating Safety Quiz, Speed Date, Video Chat (6)
In-Person / Offline Safety	Planning, attending, and following up on offline meetings	Harm during and after meetings, when the platform no longer mediates the interaction	Share My Date, Noonlight, Lyft Integration, Traveler Alert, Events, Double Date, Group Hangouts, Missed Connections, Passport Mode, Tinder U (10)
Data Access	Controlling the personal data the platform holds	Exposure or misuse of personal data	Account Deletion, Download Data, Sensitive Data Control (3)
Content Circulation Control	Limiting the reuse or spread of a user's content	Photos or messages spreading beyond the person they were shared with	Screenshot Protection, Reporting Non-Consensual Intimate Imagery, Copyright Takedown (3)
Platform Access	Entering the platform and contesting removal from it	Underage users entering an adult platform; users losing access through enforcement	Age Verification, Face Photo Requirement, Convicted Offender Exclusion, Sex Offender Registry Check, Community Guidelines Warning, Account Suspension, Account Ban, Appeals Center (8)

The first four problem areas follow the arc of dating-app use: establishing who a user is, controlling who can reach them, communicating with matches, and meeting in person. The remaining three concern the conditions of using the platform at any point: the data it holds, the spread of users' content, and who may use it at all.

Two of these areas, Identity & Authenticity and In-Person / Offline Safety, have no counterpart in Ma et al.'s [2026] framework for youth-oriented social media. Both reflect the role of dating apps in arranging meetings between strangers. Data Access and Platform Access correspond to areas in Ma et al.'s framework.

Tinder describes 56 safety features across the seven problem areas (Table 2). Contact & Visibility Control contains the most features (15; 26.8%), followed by Identity & Authenticity (11; 19.6%) and In-Person / Offline Safety (10; 17.9%). Platform Access contains 8 features (14.3%) and In-App Communication 6 (10.7%), while Data Access and Content Circulation Control contain 3 each (5.4%). These counts describe what Tinder communicates, not whether features are implemented, work as intended, or reduce risk.

The composition of each area indicates how Tinder approaches it. Identity & Authenticity centers on verification, from photo, face, and ID checks to background checks, alongside profile fields through which users present themselves. Contact & Visibility Control combines tools users apply themselves, such as unmatching, blocking, and hiding their profile, with reporting and automated detection. In-Person / Offline Safety includes protective tools for dates, such as Share My Date, Noonlight, and Traveler Alert, alongside features that structure how users meet, such as Events, Double Date, and Group Hangouts. Platform Access consists largely of enforcement: age and face-photo requirements, criminal-history screening, warnings, suspensions, bans, and appeals.

4.2 Risk Types: The 7Cs

Table 3 defines the risk types we use to map Tinder's safety features to problem areas. Content, contact, conduct, and contract risks come from the 4Cs framework [Livingstone and Stoilova 2021], circulation risk from Ma et al. [2026], and compulsivity risk from Matsumoto et al. [2026]. We add Community risk.

Table 3. The 7Cs risk taxonomy as applied to Tinder.

Risk	Definition as applied	Origin
Content	Exposure to harmful or inappropriate material	4Cs [Livingstone and Stoilova 2021]
Contact	Unwanted or harmful interaction initiated by another user	4Cs [Livingstone and Stoilova 2021]
Conduct	Risk arising from a user's own behavior, such as sending a harmful message	4Cs [Livingstone and Stoilova 2021]
Contract	Commercial or data exploitation	4Cs [Livingstone and Stoilova 2021]

Risk	Definition as applied	Origin
Circulation	Content spreading beyond its intended audience or without consent	5Cs [Ma et al. 2026]
Compulsivity	Excessive or mistimed use	6Cs [Matsumoto et al. 2026]
Community	Risk that arises once users meet offline or encounter each other within existing social networks	7Cs (this paper)

Two definitions change when the framework is applied to adults. In the 4Cs, contact risks involve interactions initiated by adults and conduct risks involve children's own behavior [Livingstone and Stoilova 2021]. On an adult platform, contact risk covers unwanted or harmful interaction initiated by any other user, and conduct risk covers risks arising from a user's own behavior.

Community risk covers what follows once users leave the platform. It includes physical safety during and after a meeting, encountering people a user already knows offline, such as a former partner or an employer, and the consequences of being identified in a particular place, such as for LGBTQ+ users traveling to countries with laws that target them [Dietzel et al. 2026; O'Brien et al. 2025]. Prior risk frameworks did not need this category, because they were developed for platforms where meeting in person is incidental. On dating apps, it is the intended outcome [O'Brien et al. 2025].

4.3 Coverage: Where Each Risk is Mitigated

Table 4. Distribution of Tinder's described safety features for each risk type across problem areas (%). Each column sums to 100%; n is the number of features addressing each risk. A feature can address more than one risk.

Problem Area	Content (n=4)	Contact (n=36)	Conduct (n=10)	Contract (n=4)	Circulation (n=3)	Compulsivity (n=1)
Identity & Authenticity	25.00	25.00	40.00	-	-	-
Contact & Visibility Control	50.00	41.67	10.00	25.00	-	100.00
In-App Communication	25.00	11.11	20.00	-	-	-
In-Person / Offline Safety	-	5.56	-	-	-	-

Problem Area	Content (n=4)	Contact (n=36)	Conduct (n=10)	Contract (n=4)	Circulation (n=3)	Compulsivity (n=1)
Data Access	-	-	-	75.00	-	-
Content Circulation Control	-	-	-	-	100.00	-
Platform Access	-	16.67	30.00	-	-	-

Content risk was addressed by four features: half in Contact & Visibility Control (50.00%), through reporting profiles and automated detection, and the remainder in Identity & Authenticity and In-App Communication (25.00% each). Tinder describes few features that address exposure to harmful material, and half of them rely on user reports or automated detection.

Contact risk was addressed by 36 features, more than any other risk. They fall mostly in Contact & Visibility Control (41.67%) and Identity & Authenticity (25.00%), followed by Platform Access (16.67%), In-App Communication (11.11%), and In-Person / Offline Safety (5.56%). Contact risk is addressed in five of the seven problem areas, suggesting that Tinder treats unwanted contact as a concern throughout use, addressed primarily through users' control over who can see and reach them and through verification of who other users are.

Conduct risk was addressed mostly in Identity & Authenticity (40.00%), followed by Platform Access (30.00%), In-App Communication (20.00%), and Contact & Visibility Control (10.00%). The features in Identity & Authenticity guide what users disclose about themselves, such as Bio Guidance and profile fields for pronouns and relationship goals, while those in Platform Access respond to violations through warnings, suspensions, and appeals. Tinder thus addresses users' own behavior both before they act, by shaping what they share, and after, through enforcement.

Contract risk was addressed mostly in Data Access (75.00%), with the remainder in Contact & Visibility Control (25.00%). Tinder frames commercial and data exploitation mainly as a matter of users' control over their personal data, through tools to download, delete, or limit the use of that data.

Circulation risk was addressed entirely in Content Circulation Control (100.00%), through screenshot protection, reporting of non-consensual intimate imagery, and copyright takedowns. Only screenshot protection acts before content spreads; the other two respond after it has been shared.

Compulsivity risk was addressed by a single feature, Disable Discovery, in Contact & Visibility Control (100.00%). Tinder presents it as a way to take a break: "Need a break from Tinder? Disable Discovery to hide your profile without deleting your account." No other described feature limits how much or when users use the app.

Community risk was addressed by 16 features, the second most of any risk. They fall mostly in In-Person / Offline Safety (56.25%), followed by Contact & Visibility Control (18.75%) and Platform Access (12.50%),

with the remainder in Identity & Authenticity and In-App Communication (6.25% each). Like contact risk, Community risk is addressed in five of the seven problem areas, but its concentration in In-Person / Offline Safety, through features such as Share My Date, Noonlight, and Traveler Alert, suggests that Tinder addresses offline and social-network risk primarily around the meeting itself.

Overall, counting each feature once for each risk it addresses, Tinder's described features most often address contact risk (36 features; 48.6%) and Community risk (16; 21.6%), followed by conduct risk (10; 13.5%), content and contract risk (4 each; 5.4%), circulation risk (3; 4.1%), and compulsivity risk (1; 1.4%). Two problem areas each address a single risk type: Data Access addresses only contract risk, and Content Circulation Control only circulation risk.

4.4 How Tinder Communicates About Its Safety Features (RQ2)

Our thematic analysis identified three patterns in how Tinder describes its safety features. A fourth pattern, in which paid features appear to distribute protection unevenly between paying and non-paying users, surfaced in a few passages but was not reviewed systematically; we return to it in Section 6.2.

4.4.1 Effectiveness Claims Without Context

Tinder's newsroom reports effects of its safety features without the baseline, sample, measurement period, or method needed to interpret them. Announcing Are You Sure? in 2021, Tinder stated that the feature "has already reduced inappropriate language in messages sent by more than 10 percent in early testing" (newsroom, 2021). For Does This Bother You?, it reported that the feature "increased reporting of harassment by 37%" (newsroom, 2021) and that "reports of inappropriate messages increased 46 percent" (newsroom, 2021), and it repeated the 46 percent figure two years later (newsroom, 2023). Other claims give no figure at all. Members who saw the Are You Sure? prompt "were less likely to be reported for inappropriate messages over the next month" (newsroom, 2021), and a newsroom post states that, "since implementing these changes," Tinder has been able to "meaningfully increase the number of trans people who remain on Tinder despite unwarranted reports" (newsroom). Statements from Match Group's Head of Safety and Social Advocacy and from the president of RAINN, both citing the features' "early results," appear in place of evidence (newsroom, 2021).

Tinder can provide context when it chooses. Its 2026 announcement of Events reports engagement figures with a note on their source, market, population, and limits (newsroom, 2026). No comparable note accompanies its safety claims.

4.4.2 Responsibility Placed on Users, Enforcement Left Undefined

Tinder's communication places the work of safety on users while leaving its own enforcement criteria undefined. Users are asked to recognize and report violations: "You know when someone's crossed the line" (policies). Reporters are told that "selecting the most appropriate reason ensures your report reaches the right team" (help center). Verification badges come with the caveat that they are not "a replacement for using your best judgment" (help center), and Photo Enhance reminds users that they are "ultimately responsible" for what they upload (help center). After harm in person, Tinder advises users to "take care of yourself and decide what you need to heal" and, in immediate danger, to "contact your local authorities right away" (help center).

Tinder's own actions, by contrast, are described in open-ended terms. Accounts are banned or suspended "when we identify violations of our Community Guidelines or Terms of Use" (help center); Tinder "may issue a warning" (help center); "repeated violations of similar behavior can result in your account being banned" (help center); and consequences range "from a nudge to a ban" (policies). After a report, Tinder will "take appropriate action" (newsroom, 2023). None of these passages states what counts as a violation, how many violations lead to which consequence, or what appropriate action involves. This pairing of individual responsibility with discretionary, unexplained enforcement matches the carceral logic Dietzel et al. [2026] identified across dating apps.

4.4.3 Availability That Varies Without Explanation

Many of Tinder's safety features are available only in some places, on some devices, or through some partners, and Tinder rarely says why. ID + Photo Verification is "only available to some people in select test markets" (help center); World ID age verification is "available in select markets" (help center); Safety Center is "currently available in select locations" (help center); profile screenshots are "blocked for users in select regions," with different behavior on Android and iOS (help center); and blocking while unmatching is offered only "if available in your region" (help center). Face Check applies only "in regions where required" (help center), and the face-photo requirement is "rolling out to more countries" without naming them (help center). Noonlight and BrightCheck are available only in the United States (help center).

When Tinder does explain variation, the reason is legal. It notifies users who exchanged messages with banned fraudulent accounts "in compliance with Colorado law" (help center), uses app-store age information "to comply with Texas law" (help center), and does not run registered sex offender checks in Maine "due to legal restrictions" (help center). Otherwise, users must work out for themselves whether a protection described in one channel applies to them. This echoes Ma et al.'s [2026] finding that rollout and regional claims leave feature availability unclear, and O'Brien et al.'s [2025] finding that safeguards vary by region.

5 Discussion

Our study adapts the problem-area framework of Ma et al. [2026] to a dating app, extends the 6Cs risk taxonomy with Community risk, and examines how Tinder communicates about its safety features. We discuss what the resulting framework shows about dating-app safety, how it can be used to evaluate where platforms address risk, and what Tinder's communication practices mean for users.

5.1 Community Risk and the Structure of Dating Apps

Prior frameworks in this lineage were built for platforms where an offline meeting between strangers is incidental. The social media platforms Ma et al. [2026] studied center on content consumption, content creation, and mediated interaction. Dating apps are organized around the opposite premise: matching, messaging, and scheduling exist to move two users toward an in-person meeting, and O'Brien et al. [2025] distinguish dating platforms from other platform types on this basis.

Community risk captures what follows once users leave the platform: physical safety during a meeting, encounters with people from a user's existing social network, and the consequences of being identified in a particular place. Because these risks arise beyond the platform, they are also where its control is

weakest, and where safety advice tends to shift responsibility onto users [Dietzel et al. 2026; O'Brien et al. 2025]. Naming Community as its own risk type makes it possible to ask how much of a platform's safety infrastructure addresses these risks, and in which problem areas. On Tinder, Community is the second most addressed risk, after contact risk. Its features concentrate around the meeting itself, through tools such as Share My Date and Noonlight, but also appear in four other problem areas (Section 4.3).

5.2 Problem Areas for Dating Apps

The problem areas we identified differ from those Ma et al. [2026] identified for youth-oriented social media. Identity & Authenticity and In-Person / Offline Safety have no counterpart in their framework. Both reflect the online-to-offline structure of dating apps: users need to know that a match is who they appear to be before meeting, and the meeting itself is a context of risk that a social media framework does not include. Because the problem areas follow the arc of dating-app use, they allow a platform's safety features to be examined stage by stage, from establishing identity to meeting in person.

5.3 Mapping Coverage as an Evaluation Tool

The main value of the framework is structural. Mapping each feature to a problem area and to the risks it claims to address shows where in the use of a platform each type of risk is mitigated, and where it is not. Ma et al. [2026] proposed their problem-area framework as a basis for auditing and comparing safety features across platforms; the 7Cs make the same comparison possible for dating apps. Applied to Tinder, the map shows where coverage is thin. Only four described features address content risk, and half of them rely on user reports or automated detection. A single feature addresses compulsivity risk. Tinder's help center states that regular use helps users "be more front and center, see more profiles and make more matches," while the only feature addressing compulsive use is a toggle for taking a break. A coverage map makes such gaps visible in a way that a list of features does not.

The resulting coverage map has several uses. Researchers could use it to select features for testing, as in audits of social media safety features [Béjar et al. 2025; Matsumoto et al. 2026]. Regulators could use it to ask whether a platform's safety communication addresses each type of risk in the contexts where it arises. Platforms could use it to describe their safety features in terms of the problem areas and risks they address, which would help users understand what protection a feature offers and when it applies.

5.4 Communication as Part of Tinder's Safety Infrastructure

The three patterns in Section 4.4 parallel those Ma et al. [2026] found for youth-oriented social media: effectiveness claims without evidence, unclear explanations of how features operate, and availability that does not match what is announced. Their recurrence on a dating platform suggests that these patterns are not specific to youth-oriented platforms. The second pattern also matches the carceral logic Dietzel et al. [2026] identified: users are asked to protect themselves, judge others, and report violations, while the platform's own decisions remain discretionary and unexplained.

These patterns point to three changes in how dating platforms communicate about safety. First, effectiveness claims should state their baseline, sample, and measurement period, as Tinder already does for engagement figures in its Events announcement. Second, platforms should describe

enforcement at the level of categories, stating which kinds of violations lead to warnings, suspensions, or bans, without disclosing thresholds that would help bad actors evade detection. Third, each description of a safety feature should state where, on which devices, and for whom the feature is available, and why availability varies. Such changes would let users understand what protection a feature offers before relying on it.

6 Limitations and Future Work

6.1 Limitations

This study has several limitations. First, we analyzed what Tinder says about its safety features, not whether those features work as described. Identity verification requirements made testing with research accounts impractical, so our findings describe claims rather than function. Second, we studied a single platform. Focusing on Tinder allowed us to compare how one platform describes the same feature across channels, but our findings may not generalize to other dating apps.

Third, problem areas were derived from the set of features Tinder describes, so the distribution of features across problem areas partly reflects how those areas were defined. Counts of features also reflect how Tinder publishes information, such as whether it describes a feature in one long help article or in several short press releases, rather than the importance or prevalence of the underlying risks. Fourth, our thematic analysis was conducted by a single researcher in two rounds of coding and should be read as preliminary. Assigning features to risk types also involved judgment, particularly for features that address more than one risk, and other researchers might draw some boundaries differently.

Finally, we analyzed only public, English-language web pages. We did not examine in-app text, onboarding flows, app store listings, advertising, regulatory filings such as transparency reports, or other public-facing media like conferences or talks. Our data also reflect Tinder's websites at the time of collection in September 2026. Because help and policy pages change without notice, Tinder's descriptions of some features may since have changed.

6.2 Future Directions

Several directions follow from this work. First, future work should examine how safety infrastructure is divided between free and paid tiers. Our dataset records whether each feature requires a subscription, which makes it possible to ask which protections are available to all users, which are limited, and which must be purchased. In our data, BrightCheck background checks offer two free searches before a fee applies, and Incognito and the settings to hide one's age and distance require a subscription. This pattern surfaced during our analysis but was not reviewed systematically. It also raises the question of whether a paid feature for one user changes the risk faced by another. The FTC's allegation that Match.com protected paying subscribers but not free users from suspected fraudulent accounts shows that such divisions can affect safety directly [CBS News 2019], and the commodification of safety features can fall unequally on marginalized users [Campaoli et al. 2025; Dietzel et al. 2026].

Second, some features reduce risk without being described as safety features. A platform may present a protective mechanism in terms of convenience, personalization, or a better experience rather than safety; Tinder's policy site, for example, describes blocking a profile as a way to keep a user's experience "a happy one." Future work can distinguish explicit safety features, which a platform presents as

protective, from implicit ones, which mitigate risk without being framed that way, and examine which risks each type addresses. This distinction would give a fuller account of a platform's safety infrastructure than its self-described safety features alone.

Third, Section 4.4.3 showed that Tinder's safety features vary in availability by region, device, and partner, and no single page documents which features exist and where they apply. Future work can systematically compare which features each channel mentions and whether their descriptions agree. Such comparisons would extend Ma et al.'s [2026] finding that rollout and regional claims leave feature availability unclear, and could inform designs such as the in-app safety dashboard Ma et al. propose, which would show users which features are active on their accounts.

Fourth, future work should audit whether dating-app safety features work as described. Béjar et al. [2025] and Matsumoto et al. [2026] tested social media safety features with scenario-based test accounts and found that many failed. Similar audits of dating apps could use our problem areas and risk mapping to decide which features to test and which risks each feature should address. Such audits must account for identity verification requirements, which make research accounts difficult to create, and may require partnerships with platforms, studies with consenting participants, or creative, novel workarounds.

Fifth, affirmative consent offers a lens for evaluating and designing dating-app safety features. Im et al. [2021] describe affirmative consent through five principles: consent is voluntary, informed, revertible, specific, and unburdensome. They argue that these principles both explain existing harms on social platforms and generate new designs. Dating apps are a natural setting for this framework, because matching, messaging, and meeting all involve negotiating consent between users. Future work can assess Tinder's features against each principle. Mutual matching, for example, requires both users to express interest before messaging, and unmatching and blocking let a user withdraw from an interaction. Sharing a match's details with friends through a feature such as Share My Date raises the question of whether the match has consented [Dietzel et al. 2026]. The same principles could guide the design of new safety features.

Sixth, future work should extend this analysis to other dating platforms, such as Bumble, Grindr, and Hinge, which O'Brien et al. [2025] also examined. These platforms differ in audience and ownership: Grindr primarily serves gay and bisexual men and trans people [Dietzel et al. 2026], while Hinge, like Tinder, is owned by Match Group [Morris 2025]. Comparing platforms would show which risks are addressed on which kinds of dating apps, and would test whether the problem areas and risk categories identified for Tinder hold elsewhere or need revision.

Finally, future work can analyze channels beyond web pages, including videos, conference presentations, and investor materials such as earnings calls and annual reports. Mapping financial figures, such as revenue and subscriber counts, onto changes in trust and safety policy over time could show whether safety commitments track user risk, regulatory pressure, or business performance. Investigative reporting suggests that these factors interact: a 2021 Match Group presentation on transparency reporting asked whether the company should publish only where required by law [The Dating Apps Reporting Project 2025].

7 Conclusion

Dating apps expose users to risks at every stage of use, and dating platforms describe a range of features meant to address them. We analyzed how Tinder describes its safety features across its help center, newsroom, and policy site. We contribute a framework for mapping a dating platform's safety features to problem areas, the contexts of use in which risks arise, and the 7Cs risk taxonomy, which extends prior frameworks with Community risk to capture risks that arise once users meet offline. Applied to Tinder, the framework shows that most described features address contact risk, that Community risk is addressed mainly around in-person meetings, and that content and compulsivity risks receive little coverage. Tinder's communication about these features presents effectiveness claims without context, places responsibility for safety on users while leaving enforcement undefined, and leaves variation in availability unexplained. We argue that dating platforms should communicate clearly which protections exist, how they work, where they apply, and which risks they address.

References

- Kath Albury, Christopher Dietzel, Tinonee Pym, et al. 2021. Not your unicorn: Trans dating app users' negotiations of personal safety and sexual health. *Health Sociology Review* 30, 1 (2021), 72–86.
- Hanan Khalid Aljasim and Douglas Zytka. 2023. Foregrounding Women's Safety in Mobile Social Matching and Dating Apps: A Participatory Design Study. *Proceedings of the ACM on Human-Computer Interaction* 7, GROUP (2023), 1–25.
- J. Banks, M. Monier, M. Reynaga, et al. 2024. From the auction block to the Tinder swipe: Black women's experiences with fetishization on dating apps. *New Media & Society* 27, 7 (2024), 4228–4245.
- Arturo Béjar, Cybersecurity for Democracy, Fairplay, Molly Rose Foundation, and ParentsSOS. 2025. *Teen Accounts, Broken Promises: How Instagram Is Failing to Protect Minors*. Technical Report. <https://fairplayforkids.org/wp-content/uploads/2025/09/Teen-Accounts-Broken-Promises-How-Instagram-is-failing-to-protect-minors.pdf>
- A. Bilz, L. A. Shepherd, and G. I. Johnson. 2023. Tainted love: A systematic literature review of online romance scam research. *Interacting with Computers* 35, 6 (2023), 773–788.
- Virginia Braun and Victoria Clarke. 2006. Using thematic analysis in psychology. *Qualitative Research in Psychology* 3, 2 (2006), 77–101. <https://doi.org/10.1191/1478088706qp063oa>
- Virginia Braun and Victoria Clarke. 2019. Reflecting on reflexive thematic analysis. *Qualitative Research in Sport, Exercise and Health* 11, 4 (2019), 589–597. <https://doi.org/10.1080/2159676X.2019.1628806>
- G. Campaioli, A. Zamperini, and M. Cecchinato. 2025. 'Now you're home': Awareness cues, rejection and post-digital safety on mobile dating apps. *New Media & Society* (2025), 1–21.
- CBS News. 2019. Government sues Match.com over fake "love interest" messages. *CBS News* (Sept. 2019). <https://www.cbsnews.com/amp/news/match-com-lawsuit-ftc-sues-match-com-over-messages-from-fake-users/>

Isha Datey and Douglas Zytke. 2024. "Just Like, Risking Your Life Here": Participatory Design of User Interactions with Risk Detection AI to Prevent Online-to-Offline Harm Through Dating Apps. *Proceedings of the ACM on Human-Computer Interaction* 8, CSCW2 (2024), 1–41.

Christopher Dietzel, Stefanie Duguay, Amy Matharu, Diana Parry, and Corey Johnson. 2026. Individualized responsibility, policing, and surveillance: Examining the carceral logic in dating app companies' approaches to safety. *Platforms & Society* 3 (2026), 1–13.
<https://doi.org/10.1177/29768624261438306>

Eric Filice, Kavishka D. Abeywickrama, Diana C. Parry, and Corey W. Johnson. 2022. Sexual violence and abuse in online dating: A scoping review. *Aggression and Violent Behavior* 67 (2022), 101781.

Rosalie Gillett. 2018. Intimate intrusions online: Studying the normalisation of abuse in dating apps. *Women's Studies International Forum* 69 (2018), 212–219.

Rosalie Gillett. 2023. 'This is not a nice safe space': Investigating women's safety work on Tinder. *Feminist Media Studies* 23, 1 (2023), 199–215.

Jane Im, Jill Dimond, Melody Berton, Una Lee, Katherine Mustelier, Mark Ackerman, and Eric Gilbert. 2021. Yes: Affirmative Consent as a Theoretical Framework for Understanding and Imagining Social Platforms. In *Proceedings of the 2021 CHI Conference on Human Factors in Computing Systems (CHI '21)*. ACM, New York, NY, USA, 18 pages. <https://doi.org/10.1145/3411764.3445778>

Sonia Livingstone and Mariya Stoilova. 2021. *The 4Cs: Classifying Online Risk to Children*. CO:RE Short Report Series on Key Topics. Leibniz-Institut für Medienforschung | Hans-Bredow-Institut (HBI), Hamburg.

Renkai Ma, Dominique Geissler, Stefan Feuerriegel, Tobias Lauinger, Damon McCoy, and Pamela J. Wisniewski. 2026. Analyzing Social Media Claims regarding Youth Online Safety Features to Identify Problem Areas and Communication Gaps. *Proceedings of the ACM on Human-Computer Interaction* 10, 2, Article CSCW005 (April 2026), 31 pages. <https://doi.org/10.1145/3788041>

Lexie Matsumoto, Arturo Béjar, Abdulraheem Arar, Damon McCoy, and Laura Edelson. 2026. *Broken, Buried, or Missing: Anatomies of Failure (and Success) of Social Media Child Safety Features*. Technical Report. Cybersafety Research Center.

Chris Morris. 2025. Match Group agrees to pay \$14 million to the FTC for misleading ads. *Fortune* (13 Aug. 2025).
<https://www.fortune.com/2025/08/13/match-group-ftc-14-million-settlement-deceptive-advertising>

Catherine R. K. O'Brien, Nuur Alifah Roslan, Steven J. Murdoch, Ruba Abu-Salma, Douglas Zytke, and Mark Warner. 2025. Online Dating Platform Safeguards and Self-Protection: How Dating Platforms Characterise, Respond to, and Safeguard Against Harms. In *Extended Abstracts of the CHI Conference on Human Factors in Computing Systems (CHI EA '25)*. ACM, New York, NY, USA, 8 pages.
<https://doi.org/10.1145/3706599.3719825>

Zahra Stardust, Rosalie Gillett, and Kath Albury. 2023. Surveillance does not equal safety: Police, data and consent on dating apps. *Crime, Media, Culture* 19, 2 (2023), 274–295.

The Dating Apps Reporting Project. 2025. Dating App Cover-Up: How Tinder, Hinge, and Their Corporate Owner Keep Rape Under Wraps. *The Markup* (13 Feb. 2025).

<https://themarkup.org/investigations/2025/02/13/dating-app-tinder-hinge-cover-up>

Emily A. Vogels and Colleen McClain. 2023. Key Findings About Online Dating in the U.S. *Pew Research Center* (2 Feb. 2023).

<https://www.pewresearch.org/short-reads/2023/02/02/key-findings-about-online-dating-in-the-u-s/>

Ge Wang, Jun Zhao, Max Van Kleek, and Nigel Shadbolt. 2021. Protection or Punishment? Relating the Design Space of Parental Control Apps and Perceptions about Them to Support Parenting for Online Safety. *Proceedings of the ACM on Human-Computer Interaction* 5, CSCW2 (2021).

<https://doi.org/10.1145/3476084>

Appendix

A. Data Collection - Detailed

Tinder Newsroom (tinderpressroom.com).

Since this site has no sitemap, the script (`scrape_tinder_pressroom_v2.py`) works off the paginated /news listing pages directly, using requests and BeautifulSoup to extract links, headlines, and dates. URLs are filtered to exclude nav and language-switcher links. Each date is matched to its link by searching that link's own local DOM container rather than assuming positional alignment across the page, which corrected an earlier bug where a duplicated carousel link would shift every subsequent date pairing by one.

Tinder policy site (policies.tinder.com).

This site's sitemap was found to omit numerous live pages (e.g., the cookie policy, accessibility statement, and several FAQ articles), so `crawl_tinder_policies.py` instead crawls recursively from a set of seed URLs, following every internal link until no new pages remain. Because the site serves the same content under three coexisting URL schemes (a bare path, an older path-based locale scheme, and a newer query-parameter scheme), each discovered URL is normalized to a single "slug identity" by stripping locale/region suffixes, retaining only the English variant of each document; an explicit `Accept-Language: en` header prevents the server from defaulting to non-English content on ambiguous bare paths. Historical, dated Privacy Policy snapshots surfaced during the crawl were retained as an archival record.

Tinder help center (help.tinder.com).

An initial general-purpose crawler, `sitemap_crawler.py`, performs a breadth-first traversal of HTML pages under a given path prefix, following internal links to build a site-wide node/edge graph. Applied to `help.tinder.com`, this approach returned a 403 error, diagnosed as a Cloudflare bot-detection challenge rather than a genuine access restriction. Since the URL structure (`/hc/en-us/...`) matched Zendesk's standard help-center convention, the script was replaced with `zendesk_help_sitemap.py`, which bypasses HTML rendering entirely by querying Zendesk's public JSON API (`categories.json`, `sections.json`, `articles.json`) directly. The script paginates through all three endpoints and computes a depth column for each article by recursively resolving `parent_section_id` relationships in the section hierarchy, since

Zendesk permits nested subsections rather than a fixed three-level structure; a category is treated as depth 1, with each level of section nesting and the containing article adding one level beyond that.

B. Extraction Fields

Table B1. Quotation table fields, adapted from Ma et al.'s four extracted items (described safety feature, affected age group, controlling entity, sub-features).

Field	Values or notes
Platform	Platform being studied
Quotation	Verbatim platform text
Described Safety Feature	Draft name(s); may list several features when a passage discusses them jointly
Affected Party	Self / Sender / Recipient / Both / Third Party / N/A; replaces affected age group
Controlling Entity	User / Platform-Automated / Platform-Manual / Third-Party
Sub-feature(s)	Features that would not exist without the parent OR are attached in description
Free/Paid Tier	Subscription requirement
Source Type	Help article / Newsroom / Policies
URL	Source traceability
Coder Notes	Sub-feature reasoning, cross-source observations, analytic commentary